

B.sop_v4.2.9 Annotation and Derived-View Format

Human-frozen character offsets are the source of truth; BIO and CWS files are reproducible derivatives.

Rule ID	Requirement	Chinese example
4.1	Authoritative record: sentence ID, original sentence, zero-based start, end-exclusive end, exact span text, and one L/K/S/T label.	sentence[start:end] == span 必须为真。
4.2	The extracted text is an exact source substring. Normalized names or ontology links, when available, are stored in a separate field.	原文 Jekins -> Jekins[S]; 不得在 span 中改写为 Jenkins。
4.3	One span carries exactly one label; spans are flat and non-overlapping. L-K-S-T is a mnemonic order, not a priority.	同句可标多个不重叠片段；同一表层跨度不得复制为两个标签。
4.4	Potential overlap is stored in a separate adjudication log, not duplicated in Gold.	字段：ID、两组候选跨度/标签、same-boundary/nested/crossing、暂定决定、最终标签、规则、仲裁人、理由。
4.5	BIO sequences are generated only after character spans are frozen and must round-trip to the same offsets.	BIO -> span -> BIO 往返后，边界和 L/K/S/T 标签必须不变。
4.6	Jieba is a boundary validator and named derived-view transform, not the human annotation authority.	检测到半词边界 -> 返回人工复核；不得自动覆盖人工 Gold。
4.7	A CWS-derived evaluation view freezes Jieba version, user dictionary, rewrite function, cap, and checksum, and names the transformed reference and prediction files.	同一派生评测中，reference 与 prediction 使用同名、同版本转换。
4.8	Model JSON output is validated against the same manual before offsets are derived; missing, extra, duplicated, crossing, or non-verbatim records fail validation.	输出 text 必须在对应 sentence 中原样出现；所有 ID 顺序与输入一致。